

Brain alignment as a usable training signal: the internal extended manuscript

MohammadErfan Jabbari

July 2, 2026

Abstract

For a person and a language model reading the same text, a simple linear map predicts the person’s brain activity from the model’s middle-layer activations. For a decade this correspondence has been used in one direction only: to measure a model’s alignment with the brain after it is trained. We ask whether it can instead be used during training. We test this in the setting of model compression: we distill a language model into a smaller student with brain alignment in the objective, at a matched budget, and ask whether the student preserves or improves that alignment, a question that remains untested. An information-budget argument implies that such a signal is far too small to act as a second teacher; at most it is a weak prior. Such a prior should help only where the text signal is weakest (small models, scarce data, aggressive compression) and wash out wherever data is abundant. **[GAP: Results summary deferred (checkpoint 2): the two established findings (the powered reality result and the no-preservation distillation result) and the downstream arc (per-individual null, averaging artifact, quality law, predictability ceiling, no practical payoff) are summarised here only once R08–R14 are consolidated; close on the payoff line then.]**

Contents

Notation and terms	2
1 Introduction	3
2 Related work	5
2.1 The classical brain-score literature	5
2.2 Brain-tuning	6
2.3 Alignment under compression	6
2.4 The empty cell	6
2.5 The measurement bar	6
3 Methods	7
3.1 The alignment score	7
3.2 What the score measures	7
3.3 Why a brain signal is only a weak prior	8
3.4 The trained–untrained gap	9
3.5 The controls	9
3.6 Data and models	9
3.7 Distillation and the brain term	10

4 Results	10
4.1 The alignment signal is real beyond low-level artifacts	10
4.2 Plain distillation does not preserve alignment by default	12
4.3 An alignment-guided objective has not yet shown a reliable held-out gain	14
4.4 No per-individual gain at matched perplexity, and the averaging artifact	14
4.5 No practical payoff	15
5 Discussion	15
6 Limitations	15
7 Conclusion	15
Checkpoint log	15

Notation and terms

Encoding model

A ridge regression from a model’s per-stimulus features to measured brain activity, scored out-of-sample. The standard brain-alignment instrument.

Voxel / ROI

A voxel is one 3D measurement location in an fMRI scan; an ROI is a predefined region (here, the five left-hemisphere language regions of the Tuckute ROI benchmark).

ΔR^2 (unique R^2)

The variance the language-model features add *on top of* the nuisance set: $R^2([Z_{\text{nuis}}, \text{LM}]) - R^2([Z_{\text{nuis}}])$. The semipartial R^2 , not the partial. The only predictive score this work reports.

Trained–untrained gap

The primary statistic: the unique R^2 of a language-trained model minus that of a same-architecture untrained network, under identical nuisance and splits. Isolates what *language training* adds over random initialisation.

Z_{nuis} (nuisance set)

The low-level nuisance variables subtracted in both arms of the gap: utterance length, position, word and phoneme rate, word duration, log word-frequency, and a static word-embedding block (eng1000). Surprisal is deliberately excluded (it is itself LM-derived).

Noise ceiling

The maximum variance any model could explain, set by the reliability of the neural measurement itself; reported here as split-half reliability of a held-out repeated story.

Permuted twin

The per-kind shuffled-target null that tests brain-*specificity*. It belongs to the lever question (whether the signal can be *moved*), not to the reality question, and so does not appear in the Results section written here.

$\mathbb{E}[Y \mid S]$

The stimulus-predictable mean response: the strong conditioning variable behind the ceiling question, as opposed to the weak low-level proxy Z_{nuis} .

$\mathcal{L}_{\text{brain}}$

The brain-alignment training loss: the term added to a language objective that rewards representations predictive of recorded brain activity.

Knowledge distillation (KD)

Training a small student model to match a larger teacher’s output distribution rather than the raw text labels. Pure logit KD is the perplexity-only form (temperature-scaled KL on the teacher’s next-token logits, no hidden-state or brain term).

Cold-init / warm-init KD

A *cold-init* student is randomly initialised before KD, so any alignment it ends with was transmitted through the distillation channel; a *warm-init* student starts from a pretrained checkpoint and is therefore a fine-tuning-drift control, not a test of what KD carries.

Headroom (Δ)

The teacher-minus-student gap in unique R^2 , $\Delta = A_T - A_S$, with A_T the teacher’s alignment and A_S the student’s. The alignment a perplexity-only objective leaves below the teacher.

Floor-anchored retention (ρ')

The fraction of the teacher’s alignment a student keeps above the untrained floor, $\rho' = (A_S - A_0)/(A_T - A_0)$, where A_0 is the untrained score: $\rho' = 0$ is the floor (destroyed), $\rho' = 1$ is the teacher (preserved). Distinct from the partial correlation $\rho_{XY \cdot Z}$ of the methods (the primed symbol is retention, never a correlation).

Data-processing inequality (DPI)

For a Markov chain $A \rightarrow B \rightarrow C$, no processing of B can increase information about A : $I(A; C) \leq I(A; B)$. Here it bounds a fixed-function compression to lose, never create, alignment relative to its teacher.

1 Introduction

Brain alignment has become a way to study what language models and the human language system share [1–3]. Linear encoding studies ask how well a model’s internal activations predict human brain responses to the same text, and the resulting score has been used to compare architectures, layers, training objectives, scale, and compression [2, 4–7]. Recent work also shows that the score is fragile: it can reflect linguistic structure, but it can also be inflated by low-level stimulus features, shallow acoustic cues, or weak evaluation splits [8–10]. Together, this literature has made brain alignment useful as a measurement. This thesis asks whether the same signal can be useful as a training signal. In particular, we study model compression: during knowledge distillation, can a small student be guided by a brain-alignment objective, and does that guidance preserve anything ordinary text-only compression loses?

Problem. We focus on model compression. Knowledge distillation trains a small *student* model to imitate a larger *teacher*, reshaping the student’s internal representations while it learns. That is where the alignment question becomes sharp: if brain alignment lives in middle-layer structure, can a brain-alignment objective shape the student while it is being distilled [2, 7]? We ask whether such a student preserves or improves alignment at a matched compute budget, and whether any gain matters beyond the alignment score itself. This case matters because compression is how small, deployable models are built, and recent brain-tuning and compression work has left this training use of brain alignment open [7, 11–13].

Hypothesis. The hypothesis is that brain alignment can act, at most, as a weak prior over text-consistent representations. The reason is a severe mismatch in scale. A language model learns from on the order of a trillion tokens of text. A language-fMRI dataset offers a few thousand noisy recordings, and a low noise ceiling caps the variance any model can explain in




Figure 1: The measurement instrument and its inversion. A ridge encoding model maps language-model activations to brain responses; held-out accuracy is reported as unique variance over the nuisance floor. The thesis asks whether a corresponding brain-alignment term can guide a distilled student.




Figure 2: The field on two axes; three cells are occupied and the fourth, where this work sits, is open in language-model compression. Examples and citations are given in Table 1.

them [1, 14, 15]. Against that imbalance, the brain term is too small to carry new task structure by itself. Its plausible role is to bias the student among solutions already compatible with the text objective, the way a weak regularizer does [16–19]. This gives a falsifiable prediction. A brain prior should help most where the text signal is weakest, in small, heavily compressed, or data-starved models, and should wash out wherever text data is abundant. Section 3 makes the argument precise.

Recent studies have already used fMRI as an optimization signal for speech and text models [11, 12, 20, 21]. The case is strengthened by causal evidence: text models pushed to be brain-*misaligned* at matched perplexity perform worse across a broad downstream suite [13]. Those studies live on the grow-and-optimize side of the field: they fine-tune the model, add adapters, or attach a brain-prediction head. Compression work has taken the complementary path. It shrinks a language model first, through quantization or pruning, and measures the alignment score afterward [7]. Alignment-guided distillation combines the two: the model is shrunk, and the alignment score enters the training objective. That is the open case in Figure 2, and it is exactly the regime where the weak-prior hypothesis says a brain signal has its best chance to matter.

Brain-alignment scores are easy to inflate. Shuffled splits leak through the temporal autocorrelation of fMRI, and low-level features such as sentence length carry much of an apparent score, so an untrained network can appear to explain brain data on its own [9, 10]. We therefore hold every number in this thesis to a fixed standard, set before measuring anything. The contributions below summarize that standard, and Section 3 defines it [2, 3, 8–10].

Contributions. This thesis makes the following contributions, scoped to the results it establishes here.

- It formulates alignment-guided distillation and locates it as the one combination the recent literature leaves open (Figure 2).
- It gives an information-budget argument that predicts where a brain prior can and cannot help, turning a vague argument for using brain data into a falsifiable prediction.
- It fixes a measurement protocol that controls for the artifacts that can fake an alignment score: unique variance over a capacity-matched nuisance floor, contiguous splits, a same-architecture untrained control, a held-out noise ceiling, and a permuted-target control for specificity, the controls the alignment-training literature has largely lacked.
- It establishes, at voxel scale and with statistical power, that the alignment signal is real beyond those artifacts [E006].
- It shows that plain perplexity-only distillation does not preserve alignment by default, while making explicit that the drop cannot yet be separated from the language-model quality the student loses [E003].
- **[GAP: The alignment-objective contribution is deferred until the R08 finding report consolidates E004; this progress draft does not claim it yet.]**

The program asks five questions in sequence, each only if the previous one holds. Is the signal real, or an artifact? Does plain distillation preserve it? Does an alignment objective move it beyond what perplexity already implies? Does any gain survive per individual at matched perplexity? And does it improve a downstream task?

The first two are settled here: the signal is real [E006], and plain distillation does not preserve it [E003]. The third question remains a marked consolidation gap until the E004 finding report is written.

2 Related work

This thesis sits in the one empty cell of a two-by-two grid: grow or shrink the model, and measure or optimize alignment (Table 1). We review the three occupied cells in turn, then the measurement standard this work adopts.

Table 1: The field on two axes: what an intervention does to the parameter budget, and how it treats brain alignment. Three cells are occupied; the fourth, shrinking a model while alignment is in the objective, is empty in the 2026 literature and is where this work sits.

	Alignment measured	Alignment optimized
Model grown	classical brain-score [2, 5, 6]	brain-tuning [11–13]
Model shrunk	compression robustness [7]	this thesis (empty in the literature)

2.1 The classical brain-score literature

The largest body of work fits encoding models to frozen, pretrained networks, then reads off how alignment varies with architecture, scale, and training. Three findings recur. Alignment lives in the structured middle layers, where syntax drives the layerwise trend and semantics is more region-specific [2]. It grows with model scale and saturates around a few billion parameters, and at a matched parameter count scale moves it more than instruction-tuning does, which is why a base model is the right teacher to align against [5]. Beyond that scale the trend can flatten or reverse, as the largest models outgrow the human language network and track formal competence more than behavior [6]. This line of work shows the signal exists and is structured, but it never intervenes through alignment, so it cannot say whether alignment is usable.

2.2 Brain-tuning

A second line puts alignment in the objective of a fine-tuned or grown model, across modalities. For speech, a voxelwise brain loss fine-tunes compact encoders and improves both alignment and several downstream tasks, with a multi-participant variant reporting further gains at higher data efficiency [11, 20]. For text, training GPT-2 and LLaMA-2 with a cosine brain loss beats text-only baselines, though the gains are mixed and the splits are shuffled with no brain-shuffled control, so they would not clear the measurement bar this work adopts [12]. The strongest evidence is causal. Training a text model to be *less* brain-aligned, at fixed perplexity, makes it worse across a suite of more than two hundred tasks, while training it to be *more* aligned improves performance on the same suite [13]. Alignment is therefore doing causal work, not just riding along with a good model. The bare inverted idea is not novel, but no study here compresses a model: every method grows or holds the parameter budget, leaving the compression question open.

2.3 Alignment under compression

The closest prior work to this thesis is the compression study of Oota et al. [7], which tracks how alignment changes as a model is compressed. They find that alignment survives only certain post-training quantizers on models of at least a few billion parameters; other quantizers, smaller models, and pruning all degrade it, and distillation is never tested. That study measures alignment under compression but never optimizes it. It reports how quantization and pruning move the score, and never makes alignment a training target. Distillation is a different and lossier operation. It refits a smaller model from scratch against the teacher, the step most likely to destroy the middle-layer structure where alignment lives. So alignment surviving quantization, a gentle reshaping of fixed weights, does not imply it survives distillation. Whether distillation preserves alignment is therefore an open empirical question: how much does it lose by default, and can a brain objective recover the rest? We develop the information-theoretic bound behind this in Section 3.

2.4 The empty cell

No prior work occupies the empty cell. Two nearby precedents, one in vision and one in audio, are sometimes cited as indirect support for a brain regularizer when data is scarce, but neither isolates a brain-specific effect. Pirlot et al. [18] regularize a vision network with neural data and improve both accuracy and adversarial robustness, yet a shuffled-label control reproduces the accuracy gain, so only the robustness benefit depends on real neural structure. Freteault et al. [19] fine-tune a small audio network on fMRI and improve most downstream tasks, but without an interaction test or a matched non-brain control, the low-data advantage is not isolated. Neither study is in the distillation regime this work targets, so both inform our design without testing the claim.

2.5 The measurement bar

A separate line of work warns that alignment scores are easy to inflate and over-read, and it sets the standard this work adopts. Shuffled splits leak through the temporal autocorrelation of fMRI, low-level features like sentence length carry much of the apparent signal, and some published scores fail basic predictive and reliability checks [9, 10]. The same headline score can reflect real semantics in one model and shallow phonology in another, so the number alone does not pin down the mechanism [8]. A differentiable, fMRI-free surrogate would go further, but the literature does not agree on its direction: local intrinsic dimension correlates with alignment in opposite ways across language and vision, and similarity-based surrogates track width and

depth as much as alignment [22–24]. We therefore leave a surrogate to future work, since real fMRI must first establish which way to push it.

3 Methods

3.1 The alignment score

We measure alignment with a ridge encoding model that maps a frozen model’s per-stimulus middle-layer activations X to recorded brain activity Y , fit on training stimuli and scored on held-out ones. For activations $X \in \mathbb{R}^{n \times d}$ over n stimuli and a v -voxel response Y , the fitted map is

$$\hat{W} = \arg \min_W \|Y - XW\|_F^2 + \lambda \|W\|_F^2, \quad (1)$$

and the score for each voxel is the correlation between its predicted and measured response. A raw score of this kind means little on its own, because a model can predict the brain for reasons that have nothing to do with language: utterance length, word position, generic word identity. We therefore report not the raw score but the variance the model adds *on top of* a low-level nuisance set Z_{nuis} , the difference of two cross-validated fits,

$$\Delta R^2 = R^2([Z_{\text{nuis}}, \text{LM}]) - R^2([Z_{\text{nuis}}]). \quad (2)$$

A positive ΔR^2 says the contextual representation explains brain activity that length, rate, and static word identity cannot [E002].

3.2 What the score measures

We use the score as a linear lower-bound proxy for an information increment: the information the model’s representation carries about brain activity beyond the low-level baseline. The underlying estimand is conditional mutual information. For three variables (X, Y, Z) ,

$$I(X; Y \mid Z) = H(Y \mid Z) - H(Y \mid X, Z) \quad (3)$$

is the dependence between X and Y that remains once Z is known [17]. Conditioning is essential here. Plain mutual information cannot distinguish the dependence of interest from a shared driver: if X and Y are both driven by some common Z , then $X \perp Y \mid Z$, so the conditional term is zero even when the unconditional term is positive. The chain rule splits the dependence into a baseline and the increment over it,

$$I(B; [Z_{\text{nuis}}, \text{LM}]) = I(B; Z_{\text{nuis}}) + I(B; \text{LM} \mid Z_{\text{nuis}}), \quad (4)$$

with B the brain response, so the information the model’s representation adds over the nuisance is exactly the conditional term $I(B; \text{LM} \mid Z_{\text{nuis}})$.

The conditional mutual information is not directly measurable from data, so we estimate it through the unique-variance score, which is its linear-Gaussian surrogate. Under joint Gaussianity, conditional mutual information is a monotone function of the squared partial correlation,

$$I(X; Y \mid Z) = -\frac{1}{2} \ln(1 - \rho_{XY \cdot Z}^2), \quad (5)$$

so the quantity is positive exactly when the partial correlation is nonzero [17, Ch. 8]. The reported score is the *semipartial* R^2 , the raw increment, which differs from the partial ρ^2 only by the factor $1/(1 - R^2([Z_{\text{nuis}}]))$. Both vanish together, so a positive semipartial score is equivalent to a positive conditional mutual information. Ridge is linear, so the score captures the linearly decodable part of that information, a population-level lower bound.

The choice of Z fixes what the claim means. Conditioning on the low-level nuisance Z_{nuis} shows the signal is real beyond low-level features, not beyond the stimulus as a whole. Conditioning instead on the full stimulus-predictable response $\mathbb{E}[Y | S]$ asks a stronger question: not whether the representation explains brain activity beyond low-level features, but whether it explains anything the stimulus itself does not already predict. That is the stimulus-predictability ceiling, and it reuses the same chain rule with $\mathbb{E}[Y | S]$ in place of Z_{nuis} as the larger conditioning set.

3.3 Why a brain signal is only a weak prior

A brain term in a language objective is, up to constants, a maximum-a-posteriori estimate. Training chooses parameters θ that minimise a loss, and adding a brain term to the usual language-model loss gives

$$\hat{\theta} = \arg \min_{\theta} \underbrace{\left[- \sum_{t=1}^N \log p_{\theta}(x_t | x_{<t}) \right]}_{-\log p(\mathcal{D}_{\text{text}} | \theta)} + \lambda \underbrace{\mathcal{L}_{\text{brain}}(\theta)}_{-\log p(\theta)}. \quad (6)$$

The first term is the negative log-likelihood of the text: next-token prediction is maximum likelihood, and minimising it maximises $p(\mathcal{D}_{\text{text}} | \theta)$. The second term is the negative log-prior: any regulariser that penalises a property of θ corresponds, up to a normalisation constant, to a prior $p(\theta) \propto \exp(-\lambda \mathcal{L}_{\text{brain}}(\theta))$ over parameters that favour that property, with λ setting its strength. Minimising their sum is therefore

$$\hat{\theta}_{\text{MAP}} = \arg \max_{\theta} [\log p(\mathcal{D}_{\text{text}} | \theta) + \log p(\theta)], \quad (7)$$

a MAP estimate in which the text is the likelihood and the brain term is the log-prior. But how far can a prior that small move a posterior dominated by the text?

A training corpus carries on the order of 10^{12} to 10^{13} bits of supervision. The brain side offers a few thousand noisy stimuli, and the information each voxel can carry is capped by its noise ceiling, a channel capacity

$$I \leq \frac{1}{2} \log_2(1 + \text{SNR}), \quad \text{SNR} = \frac{r_{\text{max}}^2}{1 - r_{\text{max}}^2}, \quad (8)$$

which puts the usable brain budget, as an order of magnitude, around 10^4 to 10^6 bits [17]. The data-processing inequality shrinks it further. Let S be the stimulus and B its noisy brain response. Under the modelling assumption that B is generated from S , the chain $\theta^* \rightarrow S \rightarrow B$ gives

$$I(\theta^*; B) \leq I(\theta^*; S). \quad (9)$$

The brain response cannot contain more information about the target parameters than the stimulus that elicited it. Such a small prior cannot teach new task structure the text does not already contain. It can only select among the many near-equivalent solutions an overparameterized model admits, which defines the role of the brain term as a regularizer rather than a source of new information.

A generalization bound can show quantitatively how and where such a weak prior can help. Modelling training as a channel from the data Z^n to the learned parameters θ , the expected generalization gap is bounded by the mutual information between them,

$$\overline{\text{gen}} \leq \sqrt{\frac{2\sigma^2 I(\theta; Z^n)}{n}}, \quad (10)$$

for σ -sub-Gaussian loss [16]. If the brain prior adds ΔI bits to $I(\theta; Z^n)$, its benefit to generalization is at most $O(\sqrt{\Delta I/n})$, and it shrinks as the task data n grows. A weak but well-targeted brain prior should therefore help most where the text signal is weakest, in small, under-trained, or aggressively compressed models, and fade where text data already dominates. Distillation is the regime where that prediction is most testable.




Figure 3: Schematic of the two computed quantities. Unique variance is the held-out R^2 increment from adding LM features to the nuisance model. The reported gap is $\Delta R^2_{\text{trained}} - \Delta R^2_{\text{untrained}}$, with both arms scored using the same nuisance set and splits. Bar lengths are illustrative.

3.4 The trained-untrained gap

We report the *gap* between a trained model and a same-architecture untrained network, both fit with the same nuisance set and the same splits, rather than the trained model’s absolute score. An absolute score is insufficient on naturalistic data, because even an untrained network can score slightly positive. The stimulus drives both the brain response and the untrained features; if the nuisance set captures only part of the stimulus, residual shared-cause dependence can remain. Subtracting the untrained arm subtracts an estimate of the component shared by architecture, tokenisation, and residual temporal structure. The gap therefore estimates what *language training* adds over random initialisation [E006].

3.5 The controls

We apply three controls uniformly. (1) *Contiguous splits*: train and test are divided in blocks (block cross-validation on the ROI benchmark, whole-story-held-out cross-validation on the voxelwise run), because shuffling time-adjacent fMRI samples leaks near-duplicate timepoints across the split and inflates the score [9][E002][E006]. (2) *A capacity-fair nuisance floor*: the wide contextual block and the wide static-embedding block are each reduced by PCA to the same rank, fit on the training fold only, so the language model does not win merely by carrying more dimensions; the few low-level scalars stay raw [9, 10][E006]. (3) *A held-out noise ceiling*: on the voxelwise run only reliable voxels are kept, selected by split-half reliability on a separate story repeated ten times and held entirely out of the cross-validation pool, so the estimate is not inflated by voxels that happen to correlate with the model by chance [14, 15][E006]. Surprisal is deliberately left out of the nuisance set, since it is itself LM-derived and subtracting it would remove part of the construct under test; log word-frequency, a static lexical scalar, is distinct and is subtracted [E006]. The permuted-target null shuffles brain responses across stimuli, breaking the stimulus-response pairing while preserving temporal structure [11, 18]. It tests *brain-specificity*, not whether the baseline alignment signal is real, and enters only where the question is whether the signal can be moved.

3.6 Data and models

We use two datasets. The Tuckute ROI benchmark uses the 1000 isolated sentences and five left-hemisphere language regions of Tuckute et al. [1]. It averages responses across training participants, so it is coarse, but its isolated-sentence design avoids the temporal-autocorrelation leakage of naturalistic time series; its reported noise ceiling is $r \approx 0.35$ [E002]. The powered measurement uses subject UTS03 of the deep-sampled naturalistic story-listening dataset of LeBel et al. [14]: 20 stories, story-level folds, and 11,442 voxels retained at split-half reliability above 0.5 on a held-out repeated story [E006]. The models are GPT-2, GPT-2-medium, and

Qwen2.5-0.5B, each read at its alignment-peak middle layer, with a same-architecture untrained network as the control in every comparison [E002][E006].

3.7 Distillation and the brain term

The compression setting is knowledge distillation: a smaller *student* learns to reproduce a larger *teacher*’s next-token distribution, rather than learning from text directly. Perplexity gives the output-quality control. Per-token cross-entropy is code length, maximum-likelihood training minimizes the KL divergence from the data distribution to the model’s, and perplexity is the exponential of that cross-entropy [17]. Both objectives are KL objectives with different references: plain language-model training projects the student onto the empirical data, while distillation projects it onto the teacher’s softened distribution. Matching *held-out perplexity* across arms holds output quality fixed while the objective varies; matching compute does not.

The distillation objective does not by itself imply preserved alignment, for two reasons. First, it constrains the *output* distribution, whereas alignment is a property of the *internal* representation, and the readout from representation to logits is many-to-one,

$$D(p_T \parallel q_S) \rightarrow 0 \not\Rightarrow Z_S \approx Y_{\text{teacher}} \not\Rightarrow I(Z_S; B \mid Z_{\text{nuis}}) = I(Y_{\text{teacher}}; B \mid Z_{\text{nuis}}), \quad (11)$$

so matching the teacher’s output need not match its representation. Second, a from-scratch student also fits the stimulus through the raw next-token targets, a route the teacher’s features do not mediate. That route breaks the Markov chain behind the data-processing bound, so the teacher’s alignment no longer upper-bounds the student’s. Whether distillation preserves alignment is therefore an empirical question, not a corollary.

The brain term enters the objective as a co-trained regression loss,

$$\mathcal{L} = \mathcal{L}_{\text{KD}} + \lambda \mathcal{L}_{\text{brain}}, \quad (12)$$

with mixing weight λ . The primary test arm is a randomly initialised student, because any alignment it retains must be learned during distillation rather than inherited from pretraining. A warm-initialised student and a no-teacher fine-tune are kept as labelled drift controls, and every arm is scored at one fixed layer within a single architecture so no layer choice biases the comparison [E003]. Brain *specificity* is tested with a permuted-target twin: the same loss, but against brain targets shuffled across stimuli. The arms share corpus, batch size, learning rate, and architecture, but they are not fully compute matched: the from-scratch arm runs about twice the optimizer steps of the warm arms, a caveat carried into the results [E003].

4 Results

4.1 The alignment signal is real beyond low-level artifacts

A trained middle layer predicts held-out language-network activity beyond low-level stimulus features, while a same-architecture untrained network does not. The reality check matters because the later intervention questions only make sense if the measured score is not a low-level artifact. The Tuckute ROI benchmark shows the effect across three language models, and the LeBel UTS03 run repeats it at voxel scale with much higher power.

The Tuckute ROI benchmark uses the five left-hemisphere language regions of Tuckute et al. [1], 1000 isolated sentences, and five contiguous folds. The trained-minus-untrained gap is positive for all three models, while the untrained control is negative at every layer and seed (Table 2). The positive signal spans a broad plateau of middle layers and disappears for random features [E002].

Table 2: Tuckute ROI benchmark (five LH language ROIs, five contiguous folds). Unique R^2 is the semipartial increment over the nuisance set; the gap is trained minus untrained. NC-norm is the gap as a fraction of the $r \approx 0.353$ noise ceiling.

Model	Best layer	Trained ΔR^2 (mean $\pm$ SD)	Untrained ΔR^2	Gap	NC-norm	Source
GPT-2 (12L)	L7	$+0.020 \pm 0.008$	-0.011	$+0.030$	0.056	[E002]
GPT-2-medium (24L)	L14	$+0.019 \pm 0.005$	-0.014	$+0.033$	0.053	[E002]
Qwen2.5-0.5B (24L)	L12	$+0.036 \pm 0.010$	-0.014	$+0.050$	0.102	[E002]

The powered voxelwise run uses subject UTS03 of LeBel et al. [14], 20 naturalistic stories, and 11,442 voxels selected at split-half reliability above 0.5 on a held-out repeated story. The main readout is the trained-minus-untrained gap: $+0.0207$ [95% CI $+0.0205$, $+0.0209$] for GPT-2 and $+0.0277$ [95% CI $+0.0274$, $+0.0280$] for Qwen2.5-0.5B (Table 3). The sign is also widespread across voxels: 95% of reliable GPT-2 voxels and 99% of reliable Qwen voxels have positive gaps. The untrained unique R^2 is -0.017 in both comparisons [E006]. The confidence intervals exclude zero. Because each trained model is compared with a same-architecture untrained control on the same stimuli, the gap points to language training rather than architecture or residual temporal structure [E006].

Table 3: Powered voxelwise run on LeBel UTS03 (20 stories, story-level folds). The gap is the trained-minus-untrained unique R^2 ; the 95% CIs quantify uncertainty on the mean gap. The percentage column reports how many reliable voxels have a positive trained-minus-untrained gap.

Model	Layer	Trained – untrained gap [95% CI]	% reliable voxels gap > 0	Source
GPT-2	L7	$+0.0207$ [95% CI $+0.0205$, $+0.0209$]	95%	[E006]
Qwen2.5-0.5B	L12	$+0.0277$ [95% CI $+0.0274$, $+0.0280$]	99%	[E006]

The controls address different ways a low-level artifact could mimic alignment. The Tuckute ROI benchmark uses isolated sentences, which removes temporal autocorrelation from naturalistic time series and makes a positive result harder to obtain [E002]. The voxelwise run uses naturalistic stimuli but holds out whole stories; test timepoints are never adjacent to training timepoints. Voxel selection uses a separate repeated story, preventing the analysis from selecting voxels because they happen to correlate with a model on the evaluation stories [E006]. Both trained and untrained scores subtract word and phoneme rate, word duration and length, log word frequency, and a static word-embedding block, the low-level feature set most likely to inflate the score [E006]. Surprisal and imageability are deliberately left unsubtracted; the nuisance baseline therefore does not remove every contextual property of the stimulus. The same-architecture untrained network supplies the stronger check: on the same stimuli and nuisance baseline, it scores negative at -0.017 . Architecture, rate, and residual temporal structure therefore cannot explain the positive trained-minus-untrained gap alone [E006].




Figure 4: Mean unique- R^2 on LeBel UTS03 reliable voxels. Trained GPT-2 L7 and Qwen2.5-0.5B L12 exceed same-architecture untrained controls; arrows show trained-minus-untrained gaps, and group labels report the fraction of reliable voxels with positive gaps [E006].

Together, the Tuckute ROI benchmark and the powered voxelwise run establish that the alignment score contains a real language-training signal. The Tuckute benchmark supplies breadth across models, and the deeply sampled UTS03 run supplies voxel-scale statistical power. Because the powered evidence comes from one subject, cross-subject generality remains open. Qwen aligns more strongly than either GPT-2 variant, while the two GPT-2 variants are nearly tied despite a threefold parameter difference. The single cross-family contrast alone cannot establish a quality ordering; Section 4.2 takes up the quality question directly [E002]. Q0 therefore gives the rest of the program a valid measurement target. The result does not show that the target can be moved during training or that moving it improves a model.

4.2 Plain distillation does not preserve alignment by default

Knowledge distillation (KD) trains a smaller student to match a larger teacher’s output distribution. E003 asks whether that ordinary perplexity-only objective keeps the teacher’s brain alignment. The from-scratch logit-KD student lands far below its teacher, with a gap of +0.0181 ($p < 0.001$), and retains only 0.37 [0.14, 0.58] of the teacher’s alignment above an untrained floor [E003]. The cold-init drop rules out automatic preservation and gives the later alignment-guided objective real headroom. That same E003 comparison also limits the claim: among the five non-teacher arms, alignment tracks log-perplexity (Pearson -0.88), so the retention gradient could reflect language-model quality rather than the distillation objective itself [E003][L011].

The headroom is $\Delta = A_T - A_S$, where A_T is the teacher’s alignment and A_S is the student’s alignment. Both quantities use the same measurement protocol as Q0: unique R^2 on the Tuckute ROI benchmark, contiguous folds, and the capacity-fair nuisance baseline. KD has to be measured directly because it follows a different path from fixed-function compression. For deterministic post-processing of a fixed teacher, such as the quantization and pruning studied by Oota et al. [7], the student’s features depend on the stimulus through the teacher’s features. The data-processing inequality then bounds $I(Z_{\text{student}}; B) \leq I(Y_{\text{teacher}}; B)$ along that teacher-to-student path. A from-scratch KD student has another route: gradient descent on the teacher’s




Figure 5: Plain KD leaves alignment headroom, with the loss entangled with model quality. (A) Floor-anchored retention ρ' falls along the rebuild ordering, from teacher-level pretrained references toward the untrained floor; the distilgpt2 retention follows the E003 lineage-teacher denominator. Error bars are bootstrap intervals where recorded. (B) Unique R^2 co-varies with held-out perplexity; the teacher is plotted for context, and the line is the recorded fit over the five non-teacher arms with Pearson coefficient -0.88 [E003].

logits still exposes the student to the training text. The fixed-function premise no longer holds, so preservation becomes an empirical question [E003].

The cold-init student is the key test. A warm-started student already inherits teacher-level alignment from pretraining, which Q0 measured directly. High retention after a few KD epochs would mainly measure fine-tuning drift from that starting point [E003]. The cold-init GPT-2 begins from random weights and is trained on the teacher’s logits, so any surviving alignment must pass through the distillation run. The warm-init KD arm and the no-teacher plain-LM finetune remain as drift controls. All comparison arms use the 12-layer GPT-2 architecture and one fixed layer, so cross-architecture layer choice cannot bias the comparison [E003].

Table 4: Knowledge-distillation arms on the Tuckute ROI benchmark at the fixed alignment layer [E003]. The ΔR^2 column reports raw unique R^2 ; E003 also reports NC-normalised values using noise ceiling 0.353. Floor-anchored retention $\rho' = (A_S - A_0)/(A_T - A_0)$ is 0 at the untrained floor and 1 at the teacher; for distilgpt2, ρ' and Δ use GPT-2, its lineage teacher. $p(\text{no-drop})$ is the bootstrap probability of no drop from the teacher.

Arm	Init / objective	ΔR^2	ρ' [boot. CI]	Δ vs teacher	$p(\text{no-drop})$	ppl	Source
teacher gpt2-medium	pretrained	+0.0188	n/a	n/a	n/a	74	[E003]
gpt2	pretrained (conventional)	+0.0195	1.02 [0.77, 1.27]	≈ 0	0.57	105	[E003]
lmft_warm	warm + plain-LM finetune	+0.0174	0.95 [0.81, 1.13]	+0.0016	0.265	44	[E003]
kd_warm	warm + pure logit KD	+0.0144	0.84 [0.67, 1.04]	+0.0046	0.061	95	[E003]
distilgpt2	real distillation (+hidden-cosine)	+0.0076	0.60 [0.42, 0.84]	+0.0119	0.001	169	[E003]
kd_cold	random + pure logit KD	+0.0005	0.37 [0.14, 0.58]	+0.0181	0.000	477	[E003]
untrained	random (floor)	-0.0102	0	n/a	n/a	$\sim 6 \times 10^4$	[E003]

The supported finding is the monotone retention gradient (Table 4). Conventional GPT-2 sits at the teacher’s level ($\rho' \approx 1.0$), warm-init KD retains 0.84, the published distilgpt2 retains 0.60 relative to GPT-2, the cold-init student retains 0.37, and the untrained floor is 0 [E003]. The reference arms therefore anchor both ends of the scale: off-the-shelf GPT-2 marks teacher-level alignment, and the random model marks the floor. Two runs scoring the same references on different GPUs reproduced the reference scores to about 0.0003 [E003]. The ordering holds at every PCA rank tested ($n_{\text{pca}} \in \{25, 50, 100\}$). The absolute magnitudes are less stable: the two quality-entangled arms change sign in unique R^2 across ranks, so the ordering is the durable finding and the loss fractions are rank-dependent [E003].

The quality caveat is important. Alignment co-varies with log-perplexity across the same non-teacher arms (Pearson -0.88). In that fit, warm-init KD lies essentially on the regression line, while conventional GPT-2 is the only arm more aligned than its perplexity predicts.

Two contrasts would separate objective-specific loss from a quality explanation, and both are marginal. Warm-init KD is less aligned than GPT-2 despite lower perplexity, with bootstrap $P(\text{gpt2} \leq \text{kd_warm}) = 0.092$. Among the warm-start controls, the no-teacher finetune beats warm-init KD on both perplexity and alignment, with $P(\text{lmft} \leq \text{kd_warm}) = 0.085$. Both contrasts lean on a single unreplicated point and, under a paired-fold test, on one fold carrying about 44% of the signal [E003][L011]. The cold-init arm also remains under-trained: its held-out perplexity is 477, far above the converged arms' 44 to 95. Part of its near-floor alignment may therefore reflect a weaker language model rather than the distillation objective itself. E003 predeclared that convergence guard [E003]. The arms were matched on corpus, batch size, learning rate, and architecture, but the from-scratch arm ran about twice the optimizer steps of the warm arms. The step imbalance deepens the under-training caveat instead of removing it [E003].

Plain perplexity-only distillation therefore leaves genuine headroom below the teacher. E003 does not decide whether an alignment-guided objective can recover that headroom beyond what perplexity already predicts. The evaluation is an ROI benchmark ($\text{NC} \approx 0.35$), not a matched-perplexity or fully optimizer-step-matched causal test. The Tuckute ROI benchmark is also underpowered for a ~ 0.005 objective-specific gap. The matched-perplexity question therefore moves to the powered voxelwise substrate [E003][L011].

4.3 An alignment-guided objective has not yet shown a reliable held-out gain

[GAP: Results Q2 (source E004): this subsection is deferred until its finding report is consolidated. The recorded result is that optimising $\mathcal{L}_{\text{brain}}$ on the strong aligner shows only a small apparent held-out alignment movement; the effect is entangled with perplexity, and at the honest fold-level inference unit the interval includes zero. This supports no demonstrated reliable gain yet, while not proving the gain is zero.]




Figure 6: Planned Q2 result figure. The final version will render the E004 held-out alignment change at the fold-level inference unit and mark the dominant fold. [GAP: Render from the E004 record and finalise the caption after the R08 finding report is consolidated.]

4.4 No per-individual gain at matched perplexity, and the averaging artifact

[GAP: Results Q3 awaits finding-reports R09–R13 (sources E005, E008, E011, E013, E017, E019, E020): the per-individual null, the cross-subject-averaging artifact, the method-general failure, the quality law, the stimulus-predictability ceiling, and the external reproduction. The parked draft reports/_pending-Q3 reclaims R09.]

4.5 No practical payoff

[GAP: Results Q4/A3 awaits finding-report R14 (source E009): the bounded OOD-perplexity null and its null-by-construction reading.]

5 Discussion

[GAP: Discussion awaits the consolidated Results arc (R07–R14). The intended spine: the positive contribution is a measurement-validity finding plus an artifact-resistant protocol, not a bare null; cross-subject target averaging manufactures apparent brain-specificity, and the matched-perplexity, per-kind permuted-twin, per-subject protocol is the control the alignment-training literature lacks. Write only once the Results sections it interprets exist.]

6 Limitations

[GAP: Limitations awaits the consolidated Results arc. The intended content, each tied to its result: the absolute unique variance is a few percent of the noise ceiling (a scale fact, not a weakness); the reality claim is "beyond the nuisance variables actually subtracted," with surprisal held out by design; the per-individual null is measured at one power level on one cohort; the stimulus-predictability ceiling is instrument-limited rather than closed. Write once the results exist.]

7 Conclusion

[GAP: Conclusion awaits the consolidated Results arc. It should restate the established positive (the alignment signal is real and powered beyond low-level artifacts) and the program's net finding once R07–R14 are consolidated, without reaching past the report layer.]

Checkpoint log

2026-06-16: Checkpoint 1, the manuscript is opened

This is the first checkpoint: the extended manuscript is created and the layers that already trace to a written finding-report are drafted. The science behind it was produced over the preceding working sessions and recorded in the experiment files; this checkpoint consolidates the part of that record that the report layer has so far narrated.

What is folded in at this checkpoint. The measurement apparatus and the one established positive result: a trained language model's middle layer carries unique variance about real language-network activity beyond a low-level nuisance regression, and a same-architecture untrained network does not [E006]. This is consolidated from the finding-report on the reality of the signal (R06), which itself rests on the Tuckute ROI benchmark [E002] and the powered voxelwise run [E006]. The front matter (Introduction, Related work) is consolidated from the background reports: the first-principles case for the brain as a weak prior (R03), the gap analysis that locates the contribution in the empty *shrink* $\times$ *optimize* cell (R04), and the synthesis of the measurement literature (R01).

What is deliberately left open. The Results sections beyond the reality of the signal (whether plain distillation preserves alignment, whether the loss is a usable lever, the per-individual null, the quality law, the predictability ceiling, and the practical-payoff question) are marked with \gap because their finding-reports (R07–R14) are not yet written. Knowledge flows down only:

a number is narrated in a report before it is consolidated here, so these sections wait on the report layer rather than reaching past it into the raw experiment records. The abstract is provisional for the same reason: it can state the question and the reality result, but not yet summarise an arc whose later findings have no narration layer.

How this checkpoint was built. The draft was written through the scientific-writing skill on the full loop, with the review pass run continuously: adversarial critique agents argued against the reasoning and the writing quality section by section, and the objections that held were addressed before the section was kept.

2026-06-22: Checkpoint 2, the distillation-preservation result and a register repair

This checkpoint does two things, neither of which changes a recorded number: it folds in the next finding-report, and it repairs the writing register the first checkpoint shipped.

What is folded in. The Q1 result, that plain perplexity-only knowledge distillation does not preserve a teacher’s brain alignment by default, is consolidated from its finding-report (R07, source [E003]) into the second Results subsection. A from-scratch logit-KD student lands a gap of +0.0181 below its teacher ($p < 0.001$, retaining 0.37 [0.14, 0.58] of the teacher’s alignment above the untrained floor), and unique R^2 falls along a monotone gradient ordered by how far each representation was rebuilt away from conventional pretraining [E003]. The stronger claim, that the distillation objective sheds alignment beyond the language-model quality it costs, is left unsupported in the text, because alignment co-varies with log-perplexity at Pearson -0.88 and the two objective-specific dissociations are only marginal [E003][L011]. The abstract and Introduction are updated to state this second result with that hedge, and their downstream-summary `\gap` markers are narrowed from R07–R14 to R08–R14. The glossary gains the terms the new section needs (knowledge distillation, cold-init/warm-init, headroom Δ , floor-anchored retention ρ' , the data-processing inequality); the retention symbol ρ' is disambiguated from the partial correlation $\rho_{XY \cdot Z}$ of the methods, a notation collision the consolidation guard flagged.

The register repair. A supervisor reading the v0.1 body flagged a storytelling register that the writing-style rules forbid (L054, D046): abstractions handed verbs of agency, the rhetorical move narrated instead of made, the repo’s internal scaffolding metaphors leaking into reader-facing text, and a reflex passive density above the threshold in the Introduction and Related work. This checkpoint rewrites those passages in the Introduction, abstract, Related work, and Results to assert the content directly, with a subject that acts and a verb that states, and varies the passive where the agent is the real topic. No claim, number, or hedge changed under the rewrite; only the voice did.

What is still open. The Discussion, Limitations, and Conclusion remain `\gap` because the results they interpret (the per-individual null, the averaging artifact, the quality law, the ceiling, the practical-payoff question) await their finding-reports R08–R14. R07’s own caveats are carried inside the Results text, as the report convention requires, rather than opened as a partial Limitations section ahead of the consolidated arc.

How this checkpoint was built. Written as its own logged `/write` session on the full loop (the ship gate D046 forbids folding a manuscript revision into another session), with the deterministic checks run on each edit and the non-skippable review pass, including the register auditor, run before handoff.

References

- [1] Greta Tuckute et al. “Driving and suppressing the human language network using large language models”. In: *Nature Human Behaviour* (2024). 1000 isolated sentences, 5 LH language ROIs; see 05-dataset-registry.

- [2] Subba Reddy Oota et al. “Joint processing of linguistic properties in brains and language models”. In: *NeurIPS* (2023). syntax drives the layer trend; alignment peaks mid-depth; see canonical/oota-2023.
- [3] Gabriele Merlin and Mariya Toneva. “Language models and brain alignment: beyond word-level and next-word prediction”. In: *Preprint* (2024). a residual survives next-word controls in IFG/AG; see canonical/merlin-2024.
- [4] Khai Loong Aw and Mariya Toneva. “Training language models to summarize narratives improves brain alignment”. In: *ICLR* (2023). objective shapes alignment, not just LM loss; see canonical/aw-2023.
- [5] Gao et al. “Scaling, not instruction-tuning, drives brain alignment”. In: *Preprint* (2024). use a base teacher, not instruct; see canonical/gao-2024.
- [6] Badr AlKhamissi et al. “Large language models can outgrow the human language network”. In: *Preprint* (2025). alignment tracks formal competence, can reverse at scale; see canonical/alkhamissi-2025.
- [7] Subba Reddy Oota et al. “Brain alignment under quantization and pruning of language models”. In: *Preprint (arXiv:2602.07547)* (2026). AWQ/SmoothQuant preserve on $\geq 3B$; GPTQ/pruning degrade; KD untested; see canonical/oota-2026.
- [8] Subba Reddy Oota et al. “Speech language models lack important brain-relevant semantics”. In: *ACL* (2024). raw scores can be hollow phonology; decompose; see canonical/oota-2024.
- [9] Nima Hadidi, Fegghi, et al. “Illusions of alignment: brain-score inflation from low-level and static features”. In: *Nature Communications* (2026). the published anti-confound protocol; see canonical/hadidi-2024. DOI: 10.1038/s41467-026-72253-7.
- [10] Jia et al. “L-PACT: prediction scores are not enough for brain alignment”. In: *Preprint* (2026). four-gate audit of frozen-encoding alignment; see canonical/jia-2026.
- [11] Moussa, Klakow, and Mariya Toneva. “Brain-tuning improves brain alignment and downstream performance of speech language models”. In: *ICLR*. L2 voxelwise loss, 90M speech LMs; see canonical/moussa-2025. 2025.
- [12] Bilgin, Leila Wehbe, et al. “Brain-informed training of text language models”. In: *ICLR*. GPT-2 124M + LLaMA-2 7B, cosine brain loss, random-shuffle splits; see canonical/bilgin-2026. 2026.
- [13] Gabriele Merlin and Mariya Toneva. “When language models lose their mind: brain (mis)alignment at matched perplexity”. In: *ICLR*. gradient-reversal de-alignment hurts 200+ tasks; see canonical/merlin-2026. 2026.
- [14] Amanda LeBel et al. “A natural language fMRI dataset for voxelwise encoding models”. In: *Scientific Data* (2023). ds003020; UTS01-03 deep-sampled; CC_norm ceiling; see canonical + 05-dataset-registry.
- [15] Alejandro Lage-Castellanos et al. “Analytical noise ceiling for fMRI encoding models”. In: *NeuroImage* (2019). formal fMRI noise-ceiling estimator; see canonical/lage-castellanos-2019.
- [16] Aolin Xu and Maxim Raginsky. “Information-theoretic analysis of generalization capability of learning algorithms”. In: *NeurIPS*. arXiv:1705.07809, Thm 1: gen bound scales as $\sqrt{2 \sigma^2 I(W; S_n)/n}$; the lecture-26 MI bound. 2017.
- [17] Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. 2nd ed. DPI, conditional MI, Gaussian closed form, rate-distortion; mapped in 06-theory-grounding. Wiley-Interscience, 2006.

- [18] Pirlot et al. “Neural-data regularization of CNNs via deep CCA to monkey V1”. In: *Preprint (arXiv:2209.02582)* (2022). accuracy gain survives a shuffled-label control; only robustness needs neural data; see canonical/pirlot-2022.
- [19] Freteault et al. “fMRI fine-tuning of a compact audio CNN improves downstream auditory tasks”. In: *Imaging Neuroscience* (2025). 12/19 HEAR tasks; low-data advantage uncontrolled; see canonical/freteault-2025.
- [20] Moussa et al. “Multi-participant brain-tuning of speech language models”. In: *NeurIPS*. per-participant LoRA, 3 training participants; see canonical/moussa-2025b. 2025.
- [21] Negi and Schwartz. “Brain-tuning improves language-model encoding of brain activity”. In: *Preprint*. reproduction target; baseline not perplexity-matched; see canonical/negi-2025. 2025.
- [22] Cheng et al. “Abstraction induces brain alignment: local intrinsic dimension and predictivity”. In: *Preprint (arXiv:2602.04081)* (2026). higher local LID -> more aligned (language); causal claim thin; see canonical/cheng-2026.
- [23] Yu et al. “Intrinsic dimension, alignment, and generalization”. In: *Preprint (arXiv:2601.22722)* (2026). lower local LID -> more aligned (vision); opposite sign to cheng-2026; see canonical/yu-2026.
- [24] Groger et al. “Aristotelian versus Platonic representational alignment”. In: *Preprint* (2026). CKA confounded by width/depth; only mutual-kNN survives null-calibration; see canonical/groger-2026.